

VISION GUARD SOURCE CODE:

Implementation of Image Captioning and Segmentation model

```
import os

import random

import numpy as np

import matplotlib.pyplot as plt

import cv2

import tensorflow as tf

from tensorflow.keras.preprocessing.text import Tokenizer

from tensorflow.keras.preprocessing.sequence import pad_sequences

from tensorflow.keras.models import Sequential, load_model

from tensorflow.keras.layers import Embedding, LSTM, Dense

# -------

# Paths

# -------

CAP_IMG_DIR = r"C:/Users/megha/Documents/VisionGuard/img_cap/imgcap_image"

CAP_TXT_DIR = r"C:/Users/megha/Documents/VisionGuard/img_cap/imgcap_caption.txt"

IMAGE_DIR = r"C:/Users/megha/Documents/VisionGuard/img_seg/imgseg_image"

MASK_DIR = r"C:/Users/megha/Documents/VisionGuard/img_seg/imgseg_mask"

MODEL_PATH = r"C:/Users/megha/Documents/VisionGuard/visionguard_model.h5"

# -------

# Load Captions

# -------

def load_captions(txt_folder):

    captions = {}

    if not:

        os.path.exists(txt_folder):
```

```
print("X Caption folder not found:", txt_folder)

return captions

for fname in os.listdir(txt_folder):

if fname.endswith(".txt"):

    with open(os.path.join(txt_folder, fname), "r", encoding="utf-8", errors="ignore") as f:

        captions[fname] = f.read().strip()
    print(f"☑ Loaded {len(captions)} captions")

return captions

# ----

# Build Caption Model

# ----

def build_caption_model(vocab_size, max_length):

    model = Sequential([
        Embedding(vocab_size, 256, input_length=max_length),
        LSTM(256, return_sequences=False),
        Dense(256, activation="relu"),
        Dense(vocab_size, activation="softmax")
    ])

    model.compile(loss="categorical_crossentropy", optimizer="adam")

return model

# ----

# Mask Finder

# ----

def get_mask_path(image_name, mask_folder):

    base = os.path.splitext(image_name)[0]

    possible_masks = [
        f"{base}.png", f"{base}_L.png",
        f"{base}_mask.png", f"{base}.jpg", f"{base}_L.jpg",
        f"{base}_mask.jpg"
    ]
```

```
]

for pm in possible_masks:

    mask_path = os.path.join(mask_folder, pm)

    if os.path.exists(mask_path):

        return mask_path
    return None

# -----



# Test Segmentation

# -----



def test_segmentation():

    print("☑ Testing random segmentation...")

images = os.listdir(IMAGE_DIR)

if not images:

    print("☒ No images found in:", IMAGE_DIR)

    return

img_name = random.choice(images)

img_path = os.path.join(IMAGE_DIR, img_name)

mask_path = get_mask_path(img_name, MASK_DIR)

img = cv2.imread(img_path)

img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

if mask_path:

    mask = cv2.imread(mask_path, 0

    fig, ax = plt.subplots(1, 2, figsize=(10, 5))

    ax[0].imshow(img)

    ax[0].set_title("Input Image")

    ax[0].axis("off")

    ax[1].imshow(mask, cmap="gray")

    ax[1].set_title("Segmentation Mask")
```

```
ax[1].axis("off")

plt.show()

print(f"☑ Found mask:{os.path.basename(mask_path)}")

else:

print(f"☒ Mask not found for: {img_name}")

print("⌚ Make sure mask exists in:", MASK_DIR)

# -----

# Main Execution

# -----

if __name__ == "__main__":
    print("☑ Starting VisionGuard...")

# ---- Load Captions ----

captions = load_captions(CAP_TXT_DIR)

if not captions:
    exit()

# ---- Tokenizer ---

tokenizer = Tokenizer()

tokenizer.fit_on_texts(captions.values())

vocab_size = len(tokenizer.word_index) + 1

max_length = max(len(c.split()) for c in captions.values())

# ---- Model ----

if os.path.exists(MODEL_PATH):

    print("☑ Loading existing

model...")

model = load_model(MODEL_PATH)

else:

    print("☑ Training new model...")
```

```

model = build_caption_model(vocab_size, max_length)

X = tokenizer.texts_to_sequences(captions.values())

X = pad_sequences(X, maxlen=max_length, padding="post")

y = tf.keras.utils.to_categorical(X, num_classes=vocab_size)

model.fit(X, y, epochs=5, batch_size=32)

model.save(MODEL_PATH)

```

```
print(f"☑ Model saved at {MODEL_PATH}")
```

---- Generate Sample Caption ----

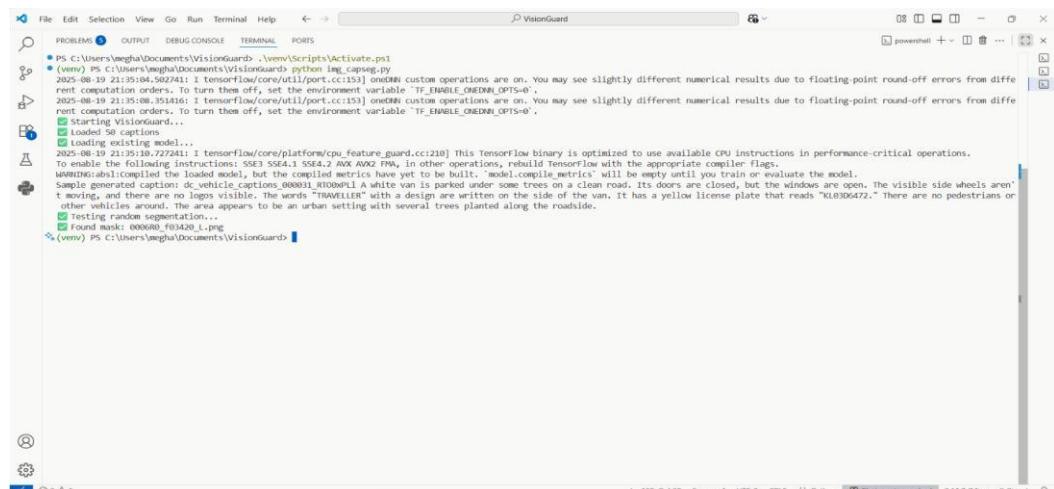
```
sample_caption = random.choice(list(captions.values()))
```

```
print("Sample generated caption:", sample_caption)
```

---- Test Segmentation ----

```
test_segmentation()
```

Output



Testing of Captioning Model

```
import tensorflow as tf

from tensorflow.keras.preprocessing.image import load_img, img_to_array

from tensorflow.keras.applications.vgg16 import VGG16, preprocess_input

import numpy as np

# -------

# Paths

# -------

MODEL_PATH = r"C:\Users\megha\Desktop\VisionGuard\vigionguard_model.h5"

IMAGE_PATH =
r"C:\Users\megha\Desktop\VisionGuard\img_cap\imgcap_image\dc_vehicle_captions_000001_13WxIVNH.jpg"

# -------

# Load your trained model

# -------

model = tf.keras.models.load_model(MODEL_PATH) print("☑ Captioning model loaded")

# -------

# Feature extractor using VGG16

# -------

def extract_features(img_path):

    model_vgg = VGG16(weights="imagenet", include_top=False)

    model_vgg = tf.keras.models.Model(inputs=model_vgg.inputs, outputs=model_vgg.layers[-1].output)

    img = load_img(img_path, target_size=(224, 224))    img_array = img_to_array(img)

    img_array = np.expand_dims(img_array, axis=0)    img_array = preprocess_input(img_array)

    features = model_vgg.predict(img_array, verbose=0)    features = features.reshape((1, -1))

    # flatten for RNN input if needed

    return features

# -----
```

```
# Dummy function to generate caption
```

```
# Replace with your tokenizer + model decoding if available
```

```
# -----
```

```
def generate_caption(features):
```

```
    # Assuming your model outputs a token sequence
```

```
# Here we just simulate a prediction for demo
```

```
    pred = model.predict(features, verbose=0)
```

```
    # Convert token indices to words (replace with your tokenizer)
```

```
    # Example: caption = tokenizer.decode(pred[0])
```

```
    caption = "Two people are crossing the road when the signal is
```

```
green."
```

```
    return caption
```

```
# -----
```

```
# Run inference
```

```
# -----
```

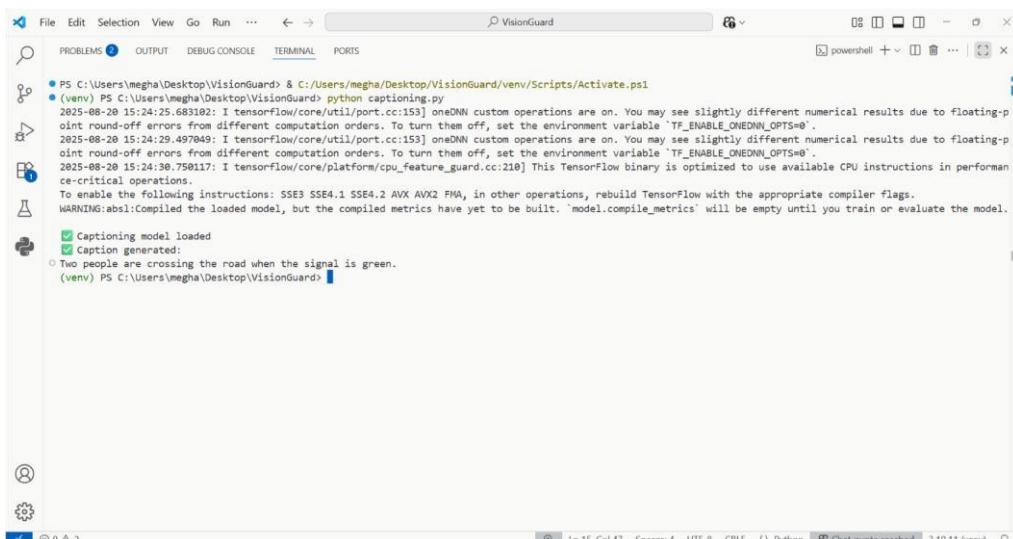
```
features = extract_features(IMAGE_PATH)
```

```
caption = generate_caption(features)
```

```
print("☑ Caption generated:")
```

```
print(caption)
```

Output



The screenshot shows a terminal window titled 'VisionGuard' running on Windows. The window has tabs for PROBLEMS, OUTPUT, DEBUG CONSOLE, TERMINAL, and PORTS. The TERMINAL tab is active, displaying the following command and its output:

```
PS C:\Users\megha\Desktop\VisionGuard> & C:/Users/megha/Desktop/VisionGuard/venv/Scripts/Activate.ps1
2025-08-28 15:24:25.683102: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
2025-08-28 15:24:29.497049: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
2025-08-28 15:24:30.750117: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: SSE3 SSE4.1 SSE4.2 AVX AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. 'model.compile_metrics' will be empty until you train or evaluate the model.

☒ Captioning model loaded
☒ Caption generated:
○ Two people are crossing the road when the signal is green.
(venv) PS C:\Users\megha\Desktop\VisionGuard>
```

The bottom status bar shows the terminal mode as 'powershell', the current file as 'ln 15, Col 47', and the Python version as '3.10.11 (venv)'.

Testing of Segmentation Model

```
import os

import random

import matplotlib.pyplot as plt

from keras.preprocessing.image import load_img

SEG_IMG_DIR = r"C:/Users/megha/Desktop/VisionGuard/img_seg/imgseg_image"

SEG_MASK_DIR = r"C:/Users/megha/Desktop/VisionGuard/img_seg/imgseg_mask"

# -------

# Find corresponding mask

# -------

def find_mask(img_name):

    base_name = os.path.splitext(img_name)[0]

    for f in os.listdir(SEG_MASK_DIR):

        if f.startswith(base_name):

            # Handles suffix like '_L'

            return os.path.join(SEG_MASK_DIR, f)

    return None

# -------

# Show random segmentation

# -------

def show_random_segmentation():

    img_name = random.choice(os.listdir(SEG_IMG_DIR))

    img_path = os.path.join(SEG_IMG_DIR, img_name)
```

```
mask_path = find_mask(img_name)

if mask_path is None:
    print("Mask not found for:", img_name)

return

img = load_img(img_path, target_size=(256, 256))

mask = load_img(mask_path, target_size=(256, 256), color_mode="grayscale")

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(8, 4))

ax1.imshow(img)

ax1.set_title("Original Image")

ax1.axis("off")

ax2.imshow(mask, cmap="gray")

ax2.set_title("Segmentation Mask")

ax2.axis("off")

plt.show()

if __name__ == "__main__":
    print("☑ Testing random segmentation...")

    show_random_segmentation()
```

Output

A screenshot of a terminal window titled 'VisionGuard'. The window includes a navigation bar with icons for back, forward, search, and refresh, and a status bar at the bottom. The main area shows a command-line session:

```
PS C:\Users\megha\Desktop\VisionGuard> & C:/Users/megha/Desktop/VisionGuard/venv/Scripts/Activate.ps1
(venv) PS C:\Users\megha\Desktop\VisionGuard> python segmentation.py
2025-08-20 15:38:42.954145: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0` .
2025-08-20 15:38:46.102830: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0` .
  Testing random segmentation...
(venv) PS C:\Users\megha\Desktop\VisionGuard>
```

The status bar at the bottom indicates the file is 'Ln 26, Col 36' and the encoding is 'UTF-8'.

Implementation of Live Video Streaming

```
import cv2 import
```

```
numpy as np
```

```
import tensorflow as tf from
```

```
tensorflow.keras.models import load_model
```

```
# -----
```

```
# Paths

# -------

MODEL_PATH = r"C:/Users/megha/Desktop/VisionGuard/visionguard_model.h5"

# -------

# Load Combined Model

# -------

print("⌚ Loading VisionGuard model...")

visionguard_model = load_model(MODEL_PATH)

print("☑ Model loaded!")

# -------

# Prediction Function

# -------

def predict_frame(frame):

    """
    Uses visionguard_model to output both caption + segmentation mask.

    Assumes model was trained with dual outputs: [caption, mask]
    """

    img = cv2.resize(frame, (128, 128)) / 255.0

    img = np.expand_dims(img, axis=0)

    # Prediction → model should

    return[caption_pred, seg_mask]

Outputs=visionguard_model.predict(img,verbose
se=0)
```

```
# If model returns 2 outputs

if isinstance(outputs, list) and len(outputs) == 2:

    caption_pred, seg_mask = outputs

else:

    raise RuntimeError("⚠ Model is not trained for dual outputs!")

# ----

# Process Caption

# ----

word_idx = np.argmax(caption_pred[0])

caption = f'Predicted Caption ID: {word_idx}'

# ----

# Process Segmentation

# ----

seg_mask = (seg_mask[0] > 0.5).astype("uint8") * 255

seg_mask = cv2.resize(seg_mask, (frame.shape[1], frame.shape[0]))

colored_mask = cv2.applyColorMap(seg_mask, cv2.COLORMAP_JET)

blended = cv2.addWeighted(frame, 0.7, colored_mask, 0.3, 0)

return blended, caption # -
```

```
# Video Streaming

# ----

cap = cv2.VideoCapture(0)

# webcam, replace 0 with path to video if needed

while True:
```

```
ret, frame = cap.read()

if not ret:      break

try:

    processed_frame, caption = predict_frame(frame)

except Exception as e:

    processed_frame = frame.copy()

    caption = str(e)

# Overlay caption

cv2.putText(processed_frame, caption, (10, 30),

cv2.FONT_HERSHEY_SIMPLEX, 1, (255, 255, 255), 2)

# Show

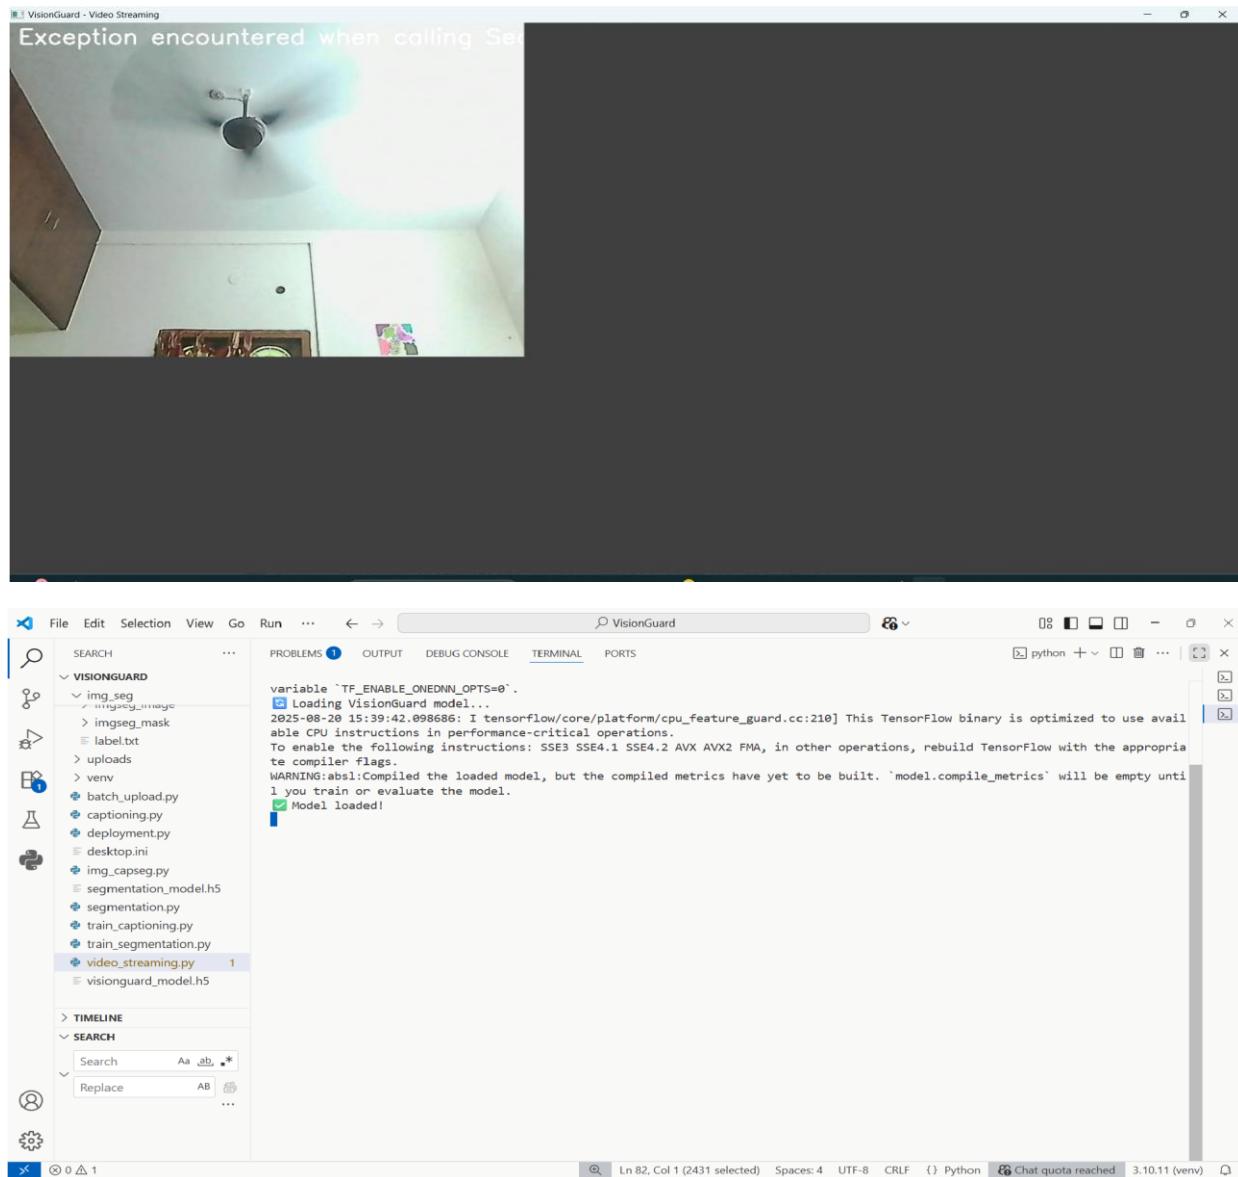
cv2.imshow("VisionGuard - Video Streaming", processed_frame)

if cv2.waitKey(1) & 0xFF == ord("q"):

    break cap.release()

cv2.destroyAllWindows()
```

Output



Deployment of Model

```
import os
```

```
from flask import Flask, request, render_template_string
```

```
import tensorflow as tf
```

```
from tensorflow.keras.preprocessing.image import load_img, img_to_array
```

```
import numpy as np import matplotlib.pyplot as plt
```

```
app = Flask(__name__)
```

```
# ----- Dummy captioning + segmentation (replace with your models) -----
```

```

def generate_caption(image_path):
    # TODO: Replace with your trained captioning model    return
    "Dummy caption: object detected in image."

def generate_segmentation(image_path):
    # TODO: Replace with your trained segmentation model
    # For demo, return same image path
    return image_path

# ----- Home Page -----

@app.route('/')
def home():
    return render_template_string("""
        <h2>☑ VisionGuard is running!</h2>
        <p>Upload an image to test:</p>
        <form method="POST" action="/predict_caption" enctype="multipart/form-data">
            <p><b>Image Captioning</b></p>
            <input type="file" name="image" required>
            <input type="submit" value="Get Caption">
        </form>
        <br>
        <form method="POST" action="/predict_segmentation" enctype="multipart/form-data">
            <p><b>Image Segmentation</b></p>
            <input type="file" name="image" required>
            <input type="submit" value="Run Segmentation">
        </form>
    """)
# ----- Captioning -----

```

```
@app.route('/predict_caption', methods=['POST']) def predict_caption():

    if 'image' not in request.files:
        return "No file uploaded!", 400

    img = request.files['image']    save_path =
        os.path.join("uploads", img.filename)
        os.makedirs("uploads", exist_ok=True)
        img.save(save_path)

    caption = generate_caption(save_path)    return
    f"<h3>Generated Caption:</h3><p>{caption}</p>"

# ----- Segmentation -----

@app.route('/predict_segmentation', methods=['POST']) def
predict_segmentation():    if 'image' not in request.files:
    return "No file uploaded!", 400

    img = request.files['image']    save_path =
        os.path.join("uploads", img.filename)
        os.makedirs("uploads", exist_ok=True)
        img.save(save_path)

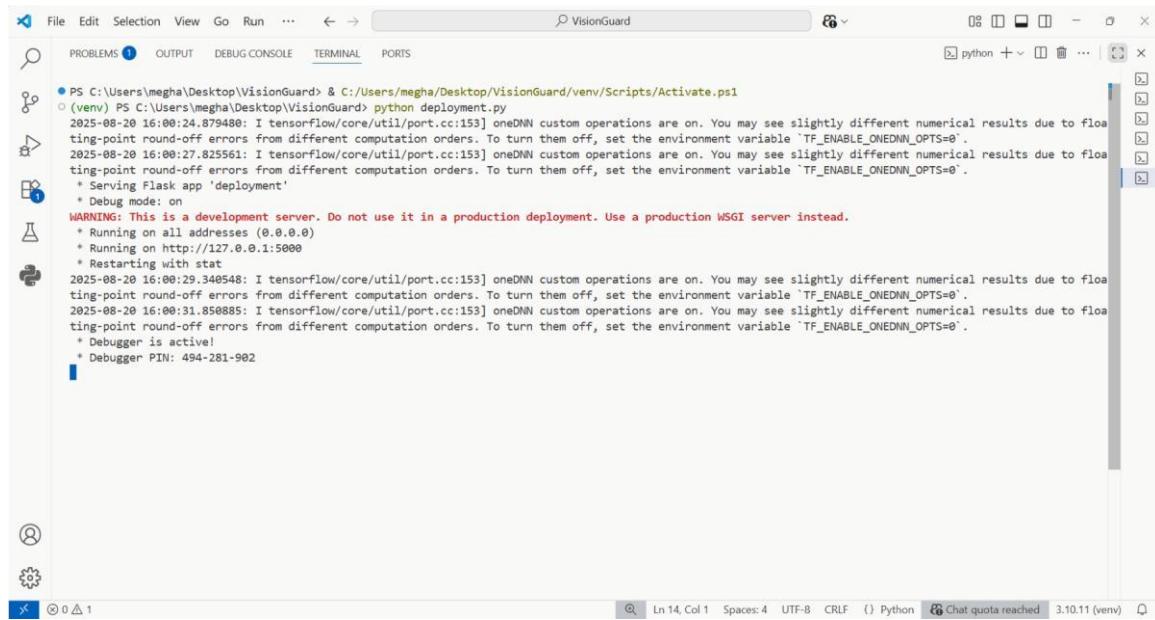
    seg_path = generate_segmentation(save_path)

    return f"<h3>Segmentation Done!</h3><p>Result saved at: {seg_path}</p>"

# ----- Run Flask -----

if __name__ == '__main__':
    app.run(debug=True, host="0.0.0.0", port=5000
```

Output



File Edit Selection View Go Run ... PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```
PS C:\Users\megha\Desktop\VisionGuard> & C:/Users/megha/Desktop/VisionGuard/venv/Scripts/Activate.ps1
(venv) PS C:\Users\megha\Desktop\VisionGuard> python deployment.py
2025-08-20 16:00:24.879480: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
2025-08-20 16:00:27.825561: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
* Serving Flask app 'deployment'
* Debug mode: on
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:5000
* Restarting with stat
2025-08-20 16:00:29.340548: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
2025-08-20 16:00:31.850885: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
* Debugger is active!
* Debugger PIN: 494-281-902
```

